

When the Window Cannot See the Question:
Token-Denominated Observation Windows Are Not
Language-Neutral

Anonymous authors
Paper under double-blind review

Abstract

054	1	Introduction
055		
056	2	Related work
057		
058	3	Setup
059		
060	4	The window is a threshold
061		
062	5	The slack must be question-sized
063		
064	6	AutoWindow: one integer, and where it stops
065		
066	7	Analysis
067		
068	8	Limitations
069		
070	9	Conclusion
071		
072		AI use statement
073		
074		Reproducibility statement
075		
076		References
077		
078	A	How last- w scoring works
079		
080	B	Measuring c and Q_{90}
081		
082	C	The scoring rule and its robustness
083		
084	D	Full window, pad, and $c+16$ grids
085		
086	E	The dose ladder in full
087		
088	F	Schema tails: the mis-sized run and the rerun
089		
090	G	A second tokenizer: Gemma-3-4b-it
091		
092	H	Scale: Qwen3-8B
093		
094	I	Frozen instructions and prompt layout
095		
096	J	Measurement transcripts and run provenance
097		
098	K	Production scope
099		
100		
101		
102		
103		
104		
105		
106		
107		