

# WHEN THE WINDOW CANNOT SEE THE QUESTION: TOKEN-DENOMINATED OBSERVATION WINDOWS ARE NOT LANGUAGE-NEUTRAL

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

Observation-window evictors such as SnapKV score the key-value cache using only the last  $w$  tokens of the prompt. We show that this integer is a language-dependent threshold. In the common `passages+question+instruction` layout, an instruction longer than  $w$  hides the question from the scorer: on Qwen3-4B the block of tokens after the question is 25 tokens in English but 107 in Bengali and 167 in Telugu, and the library default  $w=64$  costs the latter two 15 and 20 points of accuracy on the full validation pools while leaving English untouched; at the stricter constant of 32 that one serving path ships, six of our eight languages leave the window no token of the question. The failure is not linguistic. Padding an English instruction reproduces it, a tokenizer that shortens the Bengali instruction moves the threshold with it, and an English prompt carrying a JSON response schema leaves a trailing block of 166 tokens — one from Telugu’s 167 — and is blinded in exactly the same way. Recovery needs question-sized slack: a window four tokens past the instruction recovers nothing. We therefore set  $w = c + Q_{90}$  from the tokenizer, the chat template and a held-out question-length percentile. Preregistered and at matched retained cache, this removes the blind mode wherever it appears — +14 pp over the default on the full 669-item Telugu pool ( $p < 10^{-17}$ ), +13 on Bengali — closes the schema tails (non-inferior at -3 pp by interval), returns the library default where that default already meets the gate, and meets its  $\pm 3$  pp gate on all three languages at  $n=100$ ; on the full pool, Telugu instead certifies a -5.7 pp residual ( $[-7.8, -3.1]$ ). The same integers transfer unchanged to a second evictor with the same window, removing its own holes (+26/+31 pp on its own baselines). Where the gate is missed — Telugu at scale included — the residual is not a visibility failure: the question is fully in view.

## 1 INTRODUCTION

Serving a language model involves a surprising number of integers that count tokens: context limits, cache budgets, decode caps. Because tokenizers are not language-neutral (Petrov et al., 2023), each of those integers means something different in different languages, and the consequences have been documented for the budgets a user can see, such as cost, latency and usable context (Marchisio et al., 2024; Marie & Fujita, 2025). This paper is about a token-denominated integer that users cannot see. It sits inside the scoring rule of a widely used family of KV-cache eviction methods, and its default value decides whether the model is allowed to look at the question.

SnapKV (Li et al., 2024) and its descendants do not read the whole prompt when they decide what to evict. They score the cached keys against the last  $w$  tokens of the prompt — the *observation window* — and keep the highest-scoring entries. The `kvpress` implementation ships  $w=64$ ; the opt-in `RocketKV` path in `TensorRT-LLM` ships  $w=32$  — a research default and an opt-in serving path. In the retrieval-QA layout `passages + question + instruction + chat suffix`, that window is consumed from the right. Write  $c$  for the number of tokens standing after the question. If  $c > w$ , the window contains no token of the question at all, and the scorer decides what to keep without seeing what was asked (Figure 1).

On Qwen3-4B, the tokens standing after the question — a one-line QA instruction plus the chat suffix — number 25 in English, 107 in Bengali and 167 in Telugu. At the shipped default, English prompts are safe while Bengali and Telugu prompts are blind, and they lose 15 and 20 points on the full validation pools (Section 6). The mechanism, however, is not about those languages: padding the English instruction until  $c$  passes  $w$  reproduces the cliff without leaving English; a JSON response schema does the same and puts its trailing block one token from Telugu’s; and a tokenizer that encodes the Bengali instruction in 26 tokens rather than 102 moves the cliff along with it. High-fertility tokenizers are the pathway that makes a fixed instruction overflow a fixed window, not a ranking of languages by difficulty.

Clearing the trailing block turns out to be necessary but far from sufficient. A window set four tokens past  $c$  recovers nothing at all, and one that leaves an English-sized slack of sixteen tokens recovers most of the loss but not the last few points. What the scorer needs is enough *of the question*, which suggests measuring the quantity directly. We define query visibility as the share of the question inside the window,  $V = |W \cap Q|/|Q|$ , computable from the tokenizer and the template with no forward pass; Luo et al. (2026) showed that question visibility changes method rankings and left exactly such a continuous measure open.

The measurement then implies a rule that is computed rather than tuned: set the window to the trailing block plus a question-sized slack,  $\hat{w} = c + Q_{90}$ , where  $Q_{90}$  is a percentile of question length estimated on held-out questions. Both terms are measured rather than fitted to accuracy — the choice of the 90th percentile is a design decision, registered in advance, whose cost we report — no delimiter is required, and because SnapKV protects the window from within the retained budget, no extra cache is bought. Preregistered, the rule meets its  $\pm 3$  pp gate on English, Bengali and Telugu at Qwen3-4B ( $n=100$ ); closes — non-inferior by interval — the English JSON-schema tails that a mis-sized earlier version of the same rule had broken by 23 points; and returns exactly the shipped constant of 64 for a tokenizer and language where that constant is already adequate. On the full 669-item Telugu pool the gate does not hold:  $\hat{w}$  leaves  $-5.7$  pp ( $[-7.8, -3.1]$ ) beside a  $+14$  pp recovery over the default, joining the Bengali misses at Gemma-3-4b-it ( $-6$ ), Qwen3-8B and Llama-3.1-8B ( $-8$ ) — and we show the residual is not, in the main, a visibility failure: moving the question to the end of the prompt — full visibility by construction — leaves the same few points behind, and 51 of the 57 residual Telugu items have the whole question inside the window.

**Contributions.** (i) We localise a language-dependent failure of KV eviction to the relation between the observation window and the trailing block, using experiments in which language is held fixed and  $c$  moves: an English instruction pad, a change of tokenizer, and an English response schema whose trailing block lands one token from Telugu’s and is blinded the same way. (ii) We measure the dose-response, showing that the slack must be question-sized, and introduce query visibility  $V$  as an input-side, zero-inference measure of it. (iii) We propose and preregister AutoWindow,  $\hat{w} = c + Q_{90}$ , one integer per model, language and template, and report the gate, the intervals, and the misses. (iv) We diagnose the residual in the four cells where the gate is missed: the question is fully visible there, and the remaining error is of a different kind. (v) We show the integer belongs to the prompt rather than to the method: it transfers unchanged to a second evictor that inherits the same window, while the two obvious alternatives — a scorer that never reads the question, and one larger constant — miss gates the rule meets at matched ratios.

**Scope.** We study a method family rather than a shipped production default: vLLM exposes no SnapKV-style evictor, the RocketKV path is opt-in, and Mao et al. (2026) argue that research evictors rarely ship as defaults (Appendix Q). Our claim is narrower and more actionable than harm to low-resource languages in general: a token-denominated window blinds whichever prompts carry a trailing block longer than itself, and under today’s tokenizers such prompts are non-English. Which languages qualify depends on the constant: at the research default of 64 tokens, two of our eight languages are blinded; at the stricter shipped constant of 32, six are; English never is, at any window we tested. That last clause is why the number goes unaudited: at this ratio no tested window costs English anything —  $w=32$  and  $w=64$  differ by 2 pp (ns) on our items — so an English-validated benchmark cannot tell those two constants apart, and nothing in standard practice objects to whichever ships. Nor is the layout an artefact of our items: among 22 popular templates, frozen before their bodies were read, most of what ships — single-shot retrieval defaults included — is

safe, while a shipped refine template leaves exactly the library default of 64 tokens after the query and an agent scaffold leaves 221 before its first observation (Appendix P).

## 2 RELATED WORK

**Observation windows and query visibility.** Prefill eviction scores the cache from a last- $w$  slice of the prompt (Li et al., 2024). Luo et al. (2026) showed that whether the question falls inside that slice changes the *ranking* of compression methods on English long-context benchmarks — the query-aware versus query-agnostic gap is largest for SnapKV among the methods they audit — and left a continuous measure of query sensitivity open, noting that the natural candidates require a forward pass.  $V$  fills that slot from the input side: SnapKV’s accumulated attention and the estimated attention of Devoto et al. (2025) are computed from the model,  $V$  from the tokenizer and the template. Expected Attention is the informative contrast: it removes the need to see the question by estimating future queries, and never reasons about a window at all.

**How  $w$  gets chosen.** Recent eviction work treats  $w$  as a hyperparameter to set rather than as a threshold to test. IntentKV (Zhao et al., 2026) and the original SnapKV retrieval configuration use  $w=64$ ; Nawrot et al. (2026) sweep window sizes up to a few hundred tokens and select a task-average optimum. None of this work lets  $w$  fall below a trailing instruction whose length varies across otherwise matched items, which is the cell this paper occupies.

**Instruction order, protection, and delimiters.** Chen et al. (2026) document that eviction can discard an instruction in multi-instruction English prompts and propose a whitelist and a fair-eviction split; their reordering moves two instructions inside a system prompt; we move the instruction relative to the question and grow it until it overflows the window, with a remedy that needs no knowledge of prompt structure. Garcia (2026) show that protecting a prompt’s prefix and suffix outperforms scoring under a global cache cap, assuming the structurally critical tokens occupy a roughly fixed 15–30-token budget — while the Bengali and Telugu instructions occupy 102 and 162. We treat that assumption as a prediction and test it. Protection with a fixed suffix budget is, on our grid, the  $c+16$  ablation under another name — and protection with a measured budget is this paper’s rule. FINCH (Corallo & Papotti, 2024) sizes its window from a user-supplied delimiter in a context–delimiter–question layout, so  $V=1$  holds there by construction. We read that as evidence for the problem:  $w$  must adapt — but FINCH adapts via a delimiter and a question-last layout that the failure studied here does not enjoy.

**Tokenizer fertility.** Petrov et al. (2023) established that a fixed context budget is not language-neutral, and subsequent work has studied the cost, latency and usable context that follow, including under weight quantization (Marchisio et al., 2024; Marie & Fujita, 2025). Those are budgets the user can see and, in principle, plan around. The constant studied here sits inside the scorer, where it is invisible from the outside; to our knowledge no published eviction study attributes per-language damage to it.

## 3 SETUP

**Task and data.** Every item is a question, its gold passage, and same-language distractor passages packed to a fixed 8k-token prefill, with the gold passage rotating through front, middle and back positions by item index. Questions and passages come from XQuAD (Artetxe et al., 2020) for English, Chinese, Spanish, Vietnamese and Thai, and from TyDiQA-GoldP (Clark et al., 2020) for Swahili, Bengali and Telugu. The two collections are not comparable and XQuAD is parallel only within itself, so we never compare accuracy across languages. Every number in this paper is a within-language, item-paired difference between two configurations evaluated on the same 100 items.

**Models and compression.** Qwen3-4B (Qwen Team, 2025) is the primary model; Qwen3-8B is a scale slice and Gemma-3-4b-it (Gemma Team, Google DeepMind, 2025) a tokenizer contrast. We compress with SnapKV (Li et al., 2024) through the kvpress implementation at a fixed eviction ratio of 0.75 — three quarters of the prefill cache is discarded — and expose the observation window

as the treatment variable. SnapKV keeps the window inside the retained budget (Appendix A), so every comparison below is at matched retained KV. Decoding is greedy with a 384-token cap: the more common 128-token cap is itself a token-denominated constant, and Appendix L reports what it did to our own measurements.

**Scoring.** An answer counts as correct if a gold span appears, after Unicode normalization and language-aware tokenization, either inside the marked answer span or inside the first sentence of the prose. No LLM judge is used anywhere in this paper. The rule was fixed before the runs reported here and is applied offline to raw generations; a stricter, marker-only variant leaves every conclusion in the paper unchanged (Appendix C). Paired comparisons are two-sided McNemar tests, reported with their discordant counts (fixed/broken) so the reader sees what the  $p$ -value is made of. Every cell in the main text pairs 100 items; the supporting runs in Appendix O pair 200, being the same questions at two context budgets, and are marked where they appear.

**The quantities.** Prompts are instruction-last unless stated: `passages + question + instruction + chat suffix`. Write  $I$  for the token length of the trailing instruction,  $s$  for the chat-template suffix that follows it, and  $c=I+s$  for the number of tokens after the question. Write  $Q_i$  for the set of token positions holding item  $i$ 's question, so that  $|Q_i|$  is its length, and  $W$  for the set of positions in the last- $w$  window. Query visibility is the share of the question the scorer can see,  $V_i = |W \cap Q_i| / |Q_i| = \text{clamp}((w-c)/|Q_i|, 0, 1)$  in this layout, where we take the intersection as the definition and the closed form as a corollary, so that interleaved layouts do not silently break it. On Qwen3-4B the eight instructions tokenize to  $I$  between 20 (English) and 162 (Telugu) tokens; with  $s=5$  this gives  $c$  from 25 to 167 (Appendix E). Both  $c$  and  $Q$  are produced by script from the tokenizer and template actually used in the run, never copied from a table; Appendix B gives the measurement procedure and Appendix M the frozen instructions and the exact prompt layout.

**The held-out percentile.**  $Q_{90}$  is the 90th percentile of question token length on a split that excludes the 100 evaluation items: XQuAD validation after index 100, and TyDiQA-GoldP train with any evaluation question removed ( $n=1090/2387/5563$  for English, Bengali and Telugu). On Qwen3-4B this gives  $Q_{90}=18/76/80$ . It is never estimated on scored items and never tuned against accuracy.

**Software stacks.** Each comparison stays inside one model and one software stack (driver, CUDA, torch, kvpress); cells from different stacks, models or tokenizers are never pooled, even when they carry the same configuration name. The follow-up runs reported in Sections 5 and 6 re-ran cells that earlier stores already contained; across all 3,843 such generations the new text is byte-identical to the old (Appendix N), which is what licenses reading the earlier and later tables as describing one system.

## 4 THE WINDOW IS A THRESHOLD

The scorer ranks the cache against whatever occupies the tail of the prompt (Figure 1). This section establishes three facts about that tail: damage tracks the window against the tail rather than against the language; it appears in English as soon as we lengthen the tail; and it appears at a constant that is already shipped.

**A window sweep across five languages.** Table 1 reports paired damage against the uncompressed baseline for five languages at five windows, at a fixed eviction ratio of 0.75. The predictions were registered before the sweep: BLIND when  $w < c$ , SAFE when  $w > c$ , with an allowed refinement that the step sits somewhat above  $c$  rather than exactly at it.

English never moves: its trailing block is 25 tokens, so every tested window contains the whole question. Bengali is damaged in every blind cell ( $-13$  to  $-21$  pp, all  $p < .01$ ) and comes back to  $-3$  pp only at  $w=176$ , which is  $c+69$ . Telugu is damaged in all five cells, including  $w=176$  — for Telugu that window is  $c+9$  — where it is still  $-13$  pp ( $p=.019$ , uncorrected; Appendix C). Thai and Swahili, whose trailing blocks are 45 and 47 tokens, dip at the one tested window below their  $c$  and are within noise elsewhere; we do not headline Swahili, whose baseline is 56%.




Figure 1: The instruction-last layout, drawn to token scale. Shaded: the last  $w=64$  tokens, the only part of the prompt a SnapKV-style scorer reads. English leaves 25 tokens after the question, so the window covers it for all 100 English items (questions of 6–24 tokens); Bengali leaves 107 and an English prompt carrying a JSON schema leaves 166, and in both the scorer ranks the passages having seen none of the question. Question widths are item medians; the passages run on leftwards for about 8k tokens.

Table 1: Sweeping the observation window. Paired  $\Delta p$  against the uncompressed baseline, Qwen3-4B,  $n=100$  items per cell, eviction ratio 0.75. Shaded cells are those the registered rule called BLIND before the run, meaning  $w < c$ ; stars mark McNemar  $p < .05$ . Shading is a property of the design and stars of the outcome, so the table is the comparison between the two. The library default  $w=64$  sits between the third and fourth columns; its effect is Table 2.

|    | $c$ | baseline | observation window $w$ |      |      |      |      |
|----|-----|----------|------------------------|------|------|------|------|
|    |     |          | 32                     | 56   | 88   | 120  | 176  |
| en | 25  | 93       | +2                     | +1   | +1   | -1   | +1   |
| th | 45  | 85       | -6                     | -4   | 0    | -2   | -3   |
| sw | 47  | 56       | -7*                    | -4   | -2   | -4   | -3   |
| bn | 107 | 73       | -13*                   | -15* | -21* | -8   | -3   |
| te | 167 | 56       | -19*                   | -21* | -18* | -14* | -13* |

Two features matter more than the average effect. Damage follows the ordering of  $c$ , not of baseline accuracy: Swahili has the weakest baseline of the five and close to the smallest loss, while Bengali and Telugu have the longest trailing blocks and the largest losses. And clearing  $c$  is evidently not sufficient — Telugu at  $c+9$  is still down 13 points — which is what Section 5 takes up.

**The same cliff, constructed in English.** If the mechanism is trailing length rather than language, then padding an English instruction should reproduce the cliff, and widening  $w$  past the new  $c$  should remove it. We pad the English QA instruction with content-free filler to four lengths, which put  $c$  at 57, 73, 105 and 137 tokens, and sweep five windows on the same English items (Figure 2b). Accuracy at  $w=32$  is 81/86/82/78% against per-condition maxima of 92/94/93/91%.

The threshold moves with the pad: in each condition accuracy peaks at the first tested window exceeding that condition’s  $c$  and at no earlier one (the bold cells in Table 7). The sharpest case is the 96-token pad, where  $w=104$  — one token short of its  $c=105$  — scores 90% and  $w=144$  scores 93%. Plotted against the slack  $w - c$  the four conditions collapse onto one curve (Figure 2b): across the twenty cells, every window short of its condition’s trailing block scores at most 90% and every window clearing it at least 91%. Neighbouring cells sit within the noise of 100 items, so we read that separation as a description of the grid, not a test; across these curves only the number of tokens between the question and the end of the prompt changed.




Figure 2: Three sweeps of the same mechanism, each paired against its own uncompressed baseline. (a) Five languages against the window; ticks below the axis mark each language’s  $c$ , and the two languages whose trailing block exceeds most tested windows are drawn in colour. (b) English, four instruction pads, against the slack  $w - c$ : the conditions collapse onto one curve and the rule at  $w = c$  separates every cell at most 90% from every cell at least 91%. (c) Telugu, same axis and scale as (b): English is at its ceiling within about ten tokens of slack, Telugu needs eighty.

**A constant that is already shipped.** kypress ships  $w=64$ ; the opt-in RocketKV path in TensorRT-LLM ships  $w=32$ . Moving from the first constant to the second, on matched items, costs Thai 4.5 pp (2/11,  $p=.022$ ) and Swahili 6.5 pp (1/14,  $p=.001$ ;  $n=200$  each, as in Appendix O). Both languages sit safely above the threshold at  $w=64$  ( $c=45$  and  $47$ ) and below it at  $w=32$ . The failure is therefore not confined to the two scripts with the longest instructions: it reaches any language whose trailing block outgrows whichever constant is in force. Thai and Swahili would take  $\hat{w} = 91$  and  $67$ , both above their own  $c$ : the rule cannot produce this flip. A third handle points the same way. Moving the instruction to the front of the prompt puts the question inside every window without touching the press or the retained budget, and recovers Bengali by 13 pp and Telugu by 14 pp on the same items (both  $p < .02$ ; Table 2); we use it as a lever on the mechanism, not as a deployment recommendation, since a chat deployment that ends with the question already has  $V=1$ .

## 5 THE SLACK MUST BE QUESTION-SIZED

Section 4 leaves one quantity unmeasured. If the window has to reach past the trailing block before the scorer sees the question, how far past? Telugu at  $c+9$  was still 13 points down, so “clear  $c$ ” is not the answer.

**A dose ladder.** We sweep five windows at fixed offsets above the measured  $c$  on Telugu, the language with the longest trailing block and the deepest hole:  $c+4$ ,  $c+16$ ,  $c+32$ ,  $c+48$  and  $c+80$ , on 100 items with their own uncompressed baseline (Figure 2c; Table 8 in Appendix F; offsets fixed before the run,  $c$  re-measured on the run tokenizer).

Four visible tokens of question buy nothing:  $c+4$  loses 18 pp — what the same items lose at the shipped  $w=64$ , seeing none of the question. The curve then rises with the slack and reaches the baseline only at  $c+80$ . The same ladder on Bengali, a language with a different  $c$  and its own baseline, reproduces it rung for rung —  $-17$ ,  $-7$ ,  $-3$ ,  $-5$  and  $-2$  pp from  $c+4$  up to  $c+Q_{90}$  — including a  $c+4$  rung as useless there as here (Appendix F).

**Visibility as the dose.** Question length varies across items (34 to 122 tokens here), so a single window already spreads visibility across items, and each item passes through several values of  $V$  as the window widens. Pooling the five treated cells and binning by per-item  $V$ : items with  $V < 0.25$  lose 16.1 pp ( $n=118$ ,  $p < 10^{-3}$ ), while items with  $V=1$  are flat ( $+0.7$  pp,  $n=135$ ). Within items, a cluster-robust logistic regression of correctness on  $V$  across the ladder gives  $+0.68$  ( $z=3.35$ ,  $p < .001$ ), and among the items whose correctness changes along the ladder, 21 are correct at high visibility and wrong at low against 2 in the opposite direction ( $p < 10^{-4}$ ). On treated rows,  $V$  is also

Table 2: AutoWindow at a fixed eviction ratio of 0.75. Accuracy, with paired  $\Delta$ pp against each block’s own uncompressed baseline in parentheses;  $n=100$  per cell; star: McNemar  $p < .05$ . Panel (A) is Qwen3-4B on three of the eight languages measured; panel (B) is Qwen3-4B on English prompts carrying a JSON schema tail, where  $c$  is measured with the schema in place. The last column is the layout remedy — the same press and ratio with the question moved to the end of the prompt, so  $V=1$  by construction — measured against that layout’s own baseline (accuracy printed beside the difference); it was not run for panel (B), and the  $c+16$  ablation it replaces is in Appendix E. The depth re-run on the full validation pools is reported in Section 6 and Appendix D.

|                                            | window |           | accuracy ( $\Delta$ pp vs. baseline) |           |                | layout      |
|--------------------------------------------|--------|-----------|--------------------------------------|-----------|----------------|-------------|
|                                            | $c$    | $\hat{w}$ | uncompressed                         | $w=64$    | at $\hat{w}$   | $\Delta$ pp |
| <i>(A) languages</i>                       |        |           |                                      |           |                |             |
| en                                         | 25     | 43        | 93                                   | 93 (+0)   | <b>95</b> (+2) | 95 (−1)     |
| bn                                         | 107    | 183       | 73                                   | 57 (−16*) | <b>71</b> (−2) | 70 (−4)     |
| te                                         | 167    | 247       | 56                                   | 37 (−19*) | <b>56</b> (+0) | 51 (−6)     |
| <i>(B) English with a JSON schema tail</i> |        |           |                                      |           |                |             |
| JSON-60                                    | 72     | 90        | 94                                   | 88 (−6)   | <b>96</b> (+2) | —           |
| JSON-120                                   | 166    | 184       | 95                                   | 89 (−6*)  | <b>96</b> (+1) | —           |
| JSON-200                                   | 213    | 231       | 94                                   | 89 (−5)   | <b>96</b> (+2) | —           |

the narrowly better one-parameter summary: AIC prefers  $y \sim V$  (689.8) to  $y \sim (w-c)$  (692.1) and to  $y \sim (w-c) + |Q|$  (692.2) — a gap of two is support, not a discrimination.

The bins are not monotone throughout, and we state where. The  $0.75 \leq V < 1$  band loses 11.9 pp ( $n=59$ ,  $p=.065$ ), and the Bengali ladder breaks in the same band; both hold long-question items that recover once  $V$  reaches 1 (Appendix F). We claim the window-level dose curve and the within-item trend, not a monotone per-item bin profile.

Slack therefore has to be counted in units of the question rather than in constant tokens: a slack the size of an English question leaves the median Telugu item at  $V=0.31$ .

## 6 AUTOWINDOW: ONE INTEGER, AND WHERE IT STOPS

**The rule.** Both quantities that Sections 4 and 5 identify are available before any forward pass. The trailing block  $c$  comes from the tokenizer and the chat template; a question-sized slack can be estimated from a held-out sample of questions. We therefore set  $\hat{w} = c + Q_{90}$ : one integer per (model, language, template), computed offline and without a delimiter. A per-item window  $w_i = c + |Q_i|$  would give  $V_i=1$  by construction, but it requires locating the question inside the prompt, which is exactly the assumption we are trying not to make (Corallo & Papotti, 2024). The language-level percentile is a deliberately weaker commitment: at  $\hat{w}$  the window covers the entire question for 98 of 113 Bengali and 611 of 669 Telugu evaluation items (the full validation pools), and the remaining decile is registered leftover rather than a claim.

**The registered gate.** Panel (A) of Table 2 is the preregistered test ( $|\Delta| \leq 3$  pp against the uncompressed baseline). The formula meets it at  $n=100$ , and gains +14 and +19 pp on Bengali and Telugu against the shipped default (16 fixed / 2 broken and 21/2;  $p=.001$  and  $p<.001$ ). The English-sized ablation  $c+16$  behaves as Section 5 predicts: it recovers most of the hole and stops short of the baseline (−7 and −6 pp). Compared directly,  $c+16 \rightarrow \hat{w}$  is +6 and +5 pp (7 fixed / 1 broken on Telugu, 9/4 on Bengali; each ns alone, 16/5,  $p=.027$  pooled over the two languages’ independent items). Across the eight-language grid, languages whose trailing block already fits inside the default window stay within 3 pp of baseline bar one marginal exception (Appendix E); the gate is a statement about the languages the default window actually blinds. Re-run on the entire validation pools under the preregistered everything-there-is stopping rule (Appendix D), Bengali holds the gate (−1.8,  $n=113$ ) and Telugu does not: −5.7 pp (19 fixed / 57 broken, CI [−7.8, −3.1]), against a −20.2 pp hole at the default (−22.4, −17.3]) and a +14.5 pp recovery ( $p=7 \times 10^{-18}$ ). The original hundred items reproduce byte-identically and still sit at zero; the 569 new items sit at −6.7. The  $n=100$  gate was sampling luck, and we report the full-pool number as the headline.

**Schema tails, and a construction we got wrong first.** The same trailing-block logic applies when the tail is a response format rather than a translated instruction. An earlier run sized the window from the *unpadded* English instruction ( $w=41$ , below the very default it meant to improve on) and cost 23 pp at the 120-token schema. Measuring  $c$  with the schema in place gives 72, 166 and 213 tokens, hence  $\hat{w} = 90, 184$  and 231. Panel (B) shows the same formula closing all three tails while breaking none of the items the default window answered correctly (8/0, 7/0 and 7/0). By interval, the schema tails are also the best-supported close in the paper (Appendix D). The 120-token schema is worth a second look: its trailing block, 166 tokens, is within one token of Telugu’s 167. The same blindness, produced by a response format instead of a script, is removed by the same integer. Appendix G reports both runs side by side.

**A second tokenizer.** Gemma-3-4b-it tokenizes the same Bengali instruction to 26 tokens against Qwen’s 102, and its Bengali questions to a median of 11 tokens, so both inputs to the rule move:  $c$  is 27/32/44 and  $Q_{90}$  is 18/18/20 for English, Bengali and Telugu, giving  $\hat{w} = 45/50/64$ . For Telugu the rule returns exactly the constant that kvpress already ships, and Telugu is indeed flat there (−3 pp against its own baseline): a window rule that is worth running should also be able to say that the default is adequate. For Bengali,  $\hat{w}=50$  leaves −6 pp, which misses the same  $\pm 3$  pp criterion the 4B cells met (Appendix H gives the full sweep on this tokenizer).

**Scale, and a third family.** A Qwen3-8B slice reproduces the phenomenon and not the gate: the default costs Bengali 17 pp,  $\hat{w}$  recovers +9 and remains 8 below baseline ( $p=.008$ ); English is at ceiling and carries no information about either (Appendix I). Llama-3.1-8B-Instruct is the same test on a third tokenizer family, with both inputs re-measured on its tokenizer —  $c=125$  and  $Q_{90}=87$  for Bengali, so  $\hat{w}=212$ . The hole reproduces at −18 pp — the first interval in the paper entirely below the damage bound ( $[-21.5, -9.2]$ ). The formula recovers +10 pp and leaves −8 ( $p=.057$ ), the same residual as 8B; we do not claim it closes on either model (Appendix J).

**What the residuals are not.** Four cells now miss the closure gate: Bengali at Gemma (−6), at 8B and at Llama (−8), and Telugu on the primary model’s full pool (−5.7, certified). The diagnosis below was made after seeing the first three; the fourth, measured at  $6.7\times$  the sample, obeys it. With the question last, every window sees all of it by construction, and the same press at the same ratio still costs Bengali 4 pp and Telugu 6 pp (both ns) while removing about three quarters of the instruction-last hole (Table 2). The missed cells look the same: at Gemma the gap is a plateau from 48 to 64, at 8B all eight residual items and at Llama nine of eleven have their whole question inside the window, and on the full Telugu pool 51 of 57 do (Appendices H, I and J). Widening the window turns off the blind mode; it does not remove the ordinary cost of discarding three quarters of the cache.

Nor is the integer a property of one method: dropped unchanged into PyramidKV (Cai et al., 2024), which inherits the same shipped window, it takes Bengali from −28 to −2 and Telugu from −32 to −1 against that method’s own baselines (Appendix K).

## 7 ANALYSIS

Pooling the window sweep (3,000 generations over 500 items) and regressing correctness on a compression indicator and a blindness indicator  $1[w < c]$ , with standard errors clustered by item, gives +0.09 log-odds for compression when the question is visible (95% CI  $[-0.05, 0.23]$ ,  $p=.21$ ) and −1.06 for blindness ( $z=-9.4$ ,  $p<.001$ ). Uncompressed rows are never blind, so the interaction term equals the blindness term; we report the operational form. We cannot distinguish the compression cost from zero when the scorer sees the question — which is not showing it is zero: the interval admits a small effect either way. What it excludes is anything resembling the blindness term.

We also compared one-parameter descriptions of the compressed rows, and the comparison does not fully favour us. AIC prefers language fixed effects (2744) over the distance to the threshold ( $w-c$ ) (3046), the binary blindness indicator (3075), and the raw window  $w$  (3220). The mechanism’s prediction that  $(w-c)$  beats raw  $w$  holds. But on a five-language grid  $(w-c)$  is nearly collinear with language identity, and the sweep alone therefore cannot separate a threshold from a per-language constant. That separation comes from the experiments that hold language fixed while  $c$  moves (Sections 4 and 6) and from the dose ladder within one language (Section 5).



**Results that do not support the account.** A stuffed arithmetic-reasoning control does not reproduce the effect (Bengali  $-3$  pp, ns): the failure needs a question matched against passages, not merely a long prompt. Two scorers designed never to look at the question (Devoto et al., 2025; Oren et al., 2024) fail differently and worse, costing even English 14 and 8 pp at the ratio where no window we tested costs it anything (Appendix K). Four-bit KV quantization damages languages in an order that is not the fertility order. A capacity mechanism does exist, but in a different regime: at aggressive eviction, where the retained cache approaches the length of the gold passage, even English items with long passages fail, whereas at the ratio used throughout this paper Bengali damage is flat across gold-length and retention strata. We do not use high-dose fertility slopes as evidence about the window.

## 8 LIMITATIONS

The failure needs a layout in which something long stands between the question and the end of the prompt. Instruction-last prompts and schema-after-content prompts are the biting cases; a chat deployment that puts the question last has  $V=1$  by construction and nothing here applies to it.

Two properties of the remedy are by design rather than by oversight. A language-level percentile leaves the longest decile of questions below  $V=1$ ; and because SnapKV also protects the window it scores from, a very large  $w$  at a very aggressive eviction ratio would begin to tax the content budget. At the ratio used here that tax is small relative to the gold passage, which is why we did not need to separate scoring from protection to meet the gate; at  $r=0.9375$  we measure it instead of predicting it (Appendix K). Deploying the rule also assumes the prompt family is known:  $\hat{w}$  is one integer per trailing-block variant, so mixed traffic needs either per-family measurement or the maximum over families, and the mis-sized schema run in Appendix G is what skipping that measurement looks like. For a deployer the guidance splits: sizing the window removes the blind mode in every family we ran (+9 to +19 pp over the default), while the residual is a property of eviction at that ratio, addressed by lowering the ratio rather than by any window — both still true where eviction is heaviest (Appendix K).

Our items are constructed rather than drawn from a standard long-context suite, and this is a deliberate trade. Published suites vary question length, instruction length, and language together, which is precisely the confound this paper has to hold fixed, but the price is that our absolute accuracies are not comparable to leaderboard numbers. The headline cells now rest on the full validation pools — 669 Telugu and 113 Bengali items, the pools’ entire contents — which is enough to certify the Telugu hole ( $[-22.4, -17.3]$ ) and its residual ( $[-7.8, -3.1]$ ), while Bengali’s capped pool leaves its interval wide ( $[-8.0, +5.2]$ ). Finally, one primary 4B model carries the headline result; the 8B slice shares its tokenizer, so it is a replication in scale rather than an independent sample, and a single contrast tokenizer cannot establish how the rule behaves across the space of tokenizers.

## 9 CONCLUSION

A token-denominated observation window is not language-neutral, because trailing instructions are not. Two quantities decide the outcome and both can be computed before any inference:  $c$ , the number of tokens standing between the question and the end of the prompt, and the slack the window leaves beyond it, which has to be counted in units of the question rather than in constants: too little slack is blindness, and too much, at heavy eviction, is a tax (Appendix K). Setting  $w = c + Q_{90}$  from the tokenizer, the chat template and a held-out percentile removes the blind mode, returns the library default where that default already suffices, leaves a certified residual that is not a visibility failure, and transfers unchanged to a second evictor — the integer belongs to the prompt, not the method. The general lesson is cheaper than the specific one, and we did not escape it ourselves: our own decode cap chose which branch of the scoring rule judged each language (Appendix L). Any constant that counts tokens inherits the tokenizer; when it gates what a model may look at, it should be derived from the tokenizer and the template, not shipped as a number.

## ETHICS STATEMENT

This work uses publicly released question-answering corpora and publicly released model weights; it involves no human subjects and no personal data. The failure it documents falls unevenly on users who write in scripts that current tokenizers encode inefficiently, so we have tried to describe it in terms of the mechanism rather than in terms of language rankings, and to avoid comparisons of absolute accuracy between languages that our item construction does not support. The remedy we propose is cheap and requires no additional data collection.

## REPRODUCIBILITY STATEMENT

Configurations, seeds, prompts and the scoring rule are in the supplementary code. The two constants are produced by scripts, not copied between documents: `measure_c.py` (and `measure_c_schema.py` for schema tails) computes  $c$  from the run tokenizer and chat template, and `measure_q.py` computes  $Q_{90}$  from a held-out question split. The dose-response readout is produced by `v_trace_bins.py`, which was committed before the corresponding data existed. Eighteen generation stores back the main arms — `cliff_multi`, `cliff_en`, `autowin`, `autowin_q90`, `autowin_8b`, `cliff_gemma`, `gemma_q90`, `schema`, `schema_fix`, `v_trace`, `v_trace_bn`, `llama`, `instr_first`, `agnostic`, `ratio`, `pyramidkv`, `constant` and `depth` — and Appendix L draws on two further cap-128-era stores, `e3` and `e3.384`. Every table cell in this paper is recomputed from the raw generations in those twenty stores by `audit_paper_numbers.py`, which fails on any drift, and comparisons never cross a stack boundary (Appendix N). The predictions and kill conditions for every arm are files that ship with the code, each written before its run; for nine of the twelve arms the version-history timestamp also precedes the first generation in the corresponding store, while the other three files entered version control in a later batch. The anonymized version history ships with the supplementary material, and the de-anonymized history becomes public with the code on publication.

## AI USE STATEMENT

Large language models were used substantially and throughout this work, and this statement follows the categories of the ICLR AI policy. Research ideation, experimental design and the pre-registered predictions were developed by the author in dialogue with LLMs; every preregistration was reviewed, approved and committed by the author before the corresponding run. Experiments were executed by LLM agents operating the experiment harness against those committed protocols, whose kill conditions bound what an agent could decide alone. The harness, the analysis scripts and the figures are code written with substantial LLM assistance; the paper’s text was drafted and revised the same way. Interpretation of results was constrained by design rather than delegated: no LLM judges any output in any reported metric, every table cell is recomputed from the raw generation records by a committed audit script, and cross-checking between independent models during preparation surfaced and corrected errors, including fabricated bibliography entries and two internal miscounts, before submission. LLMs were not used to generate synthetic evaluation data; all items derive from XQuAD and TyDiQA-GoldP. The author reviewed all AI-assisted content and takes full responsibility for every claim in this paper.

## REFERENCES

- Mikel Artetxe, Sebastian Ruder, and Dani Yogatama. On the cross-lingual transferability of monolingual representations. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 4623–4637, 2020. doi: 10.18653/v1/2020.acl-main.421.
- Zefan Cai, Yichi Zhang, Bofei Gao, Yuliang Liu, Yucheng Li, Tianyu Liu, Keming Lu, Wayne Xiong, Yue Dong, Junjie Hu, and Wen Xiao. PyramidKV: Dynamic KV cache compression based on pyramidal information funneling. *arXiv preprint arXiv:2406.02069*, 2024.
- Alex Chen, Renato Geh, Aditya Grover, Guy Van Den Broeck, and Daniel Mingyi Israel. The pitfalls of KV cache compression. In *Proceedings of the 64th Annual Meeting of the Association*

- for *Computational Linguistics (Volume 1: Long Papers)*, pp. 41530–41553, 2026. doi: 10.18653/v1/2026.acl-long.1926. ACL Anthology 2026.acl-long.1926.
- Jonathan H. Clark, Eunsol Choi, Michael Collins, Dan Garrette, Tom Kwiatkowski, Vitaly Nikolaev, and Jennimaria Palomaki. TyDi QA: A benchmark for information-seeking question answering in typologically diverse languages. *Transactions of the Association for Computational Linguistics*, 8:454–470, 2020. doi: 10.1162/tac1\_a.00317.
- Giulio Corallo and Paolo Papotti. FINCH: Prompt-guided key-value cache compression for large language models. *Transactions of the Association for Computational Linguistics*, 12:1517–1532, 2024. doi: 10.1162/tac1\_a.00716. Anthology 2024.tac1-1.83.
- Alessio Devoto, Maximilian Jeblick, and Simon Jégou. Expected attention: KV cache compression by estimating attention from future queries distribution. *arXiv preprint arXiv:2510.00636*, 2025.
- Gabriel Garcia. Protection is (nearly) all you need: Structural protection dominates scoring in globally capped KV eviction. *arXiv preprint arXiv:2605.18053*, 2026.
- Gemma Team, Google DeepMind. Gemma 3 technical report, 2025. arXiv:2503.19786.
- Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, Acyr Locatelli, Hanchen Ye, Tianle Cai, Patrick Lewis, and Deming Chen. SnapKV: LLM knows what you are looking for before generation. In *Advances in Neural Information Processing Systems*, 2024.
- Daming Luo, Christy Liang, and Junyu Xuan. How query visibility changes KV-cache compression rankings: A matched-budget audit. *arXiv preprint arXiv:2607.11942*, 2026.
- Weian Mao, Yukang Chen, Wei Huang, Shuai Yang, Luozhou Wang, and Song Han. KV cache compression and its infra problems. NVIDIA Efficient AI Lab research blog, 12 June 2026, 2026. Cited against a serving-default claim.
- Kelly Marchisio, Saurabh Dash, Hongyu Chen, Dennis Aumiller, Ahmet Üstün, Sara Hooker, and Sebastian Ruder. How does quantization affect multilingual LLMs? *arXiv preprint arXiv:2407.03211*, 2024.
- Benjamin Marie and Atsushi Fujita. The uneven impact of post-training quantization in machine translation. *arXiv preprint arXiv:2508.20893*, 2025.
- Piotr Nawrot, Jianing Li, Renjie Huang, Sebastian Ruder, Kelly Marchisio, and Edoardo Ponti. The sparse frontier: Sparse attention trade-offs in transformer LLMs. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 38667–38701, 2026. doi: 10.18653/v1/2026.findings-acl.1926. ACL Anthology 2026.findings-acl.1926; arXiv:2504.17768.
- Matanel Oren, Michael Hassid, Nir Yarden, Yossi Adi, and Roy Schwartz. Transformers are multi-state RNNs. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 18724–18741, 2024.
- Aleksandar Petrov, Emanuele La Malfa, Philip Torr, and Adel Bibi. Language model tokenizers introduce unfairness between languages. In *Advances in Neural Information Processing Systems*, 2023.
- Qwen Team. Qwen3 technical report, 2025. arXiv:2505.09388.
- Liang Zhao, Xiaocheng Feng, Weihong Zhong, Lei Huang, Kun Zhu, Baoxin Wang, Dayong Wu, Guoping Hu, Ting Liu, and Bing Qin. Question tells you where the answer is: Intention-aware long-context KV cache compression. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 27153–27169, 2026. doi: 10.18653/v1/2026.acl-long.1250. ACL Anthology 2026.acl-long.1250.

## A HOW LAST- $w$ SCORING WORKS

SnapKV forms  $q_{\text{obs}} = q[:, :, -w :]$ , scores the prefix keys against that slice, keeps the highest-scoring entries under the compression ratio, and protects the window itself from eviction. The kvpress implementation defaults to  $w=64$  with a pooling kernel of size 5, and the RocketKV path skips sparse KV entirely when the sequence is shorter than its prompt budget. Because the window is retained rather than added, raising  $w$  at a fixed compression ratio does not increase the amount of cache that survives: it changes which entries are selected, not how many.

## B MEASURING $c$ AND $Q_{90}$

$c$  is produced by `measure_c.py`. A probe prompt is wrapped in the model’s chat template with `add_generation_prompt`; if the probe question’s token ids appear as a contiguous subsequence,  $c$  is the number of tokens after that span (`after_question`); otherwise  $c = I + s$ , where  $s$  is the assistant-header suffix measured on a minimal turn (`fallback_I_plus_suffix`). Qwen3-4B takes the fallback path with  $s=5$  and Gemma-3-4b-it the direct path with  $s=6$ . We verified the fallback against a direct offset-mapping measurement on real prompts: for all eight Qwen languages the two agree within one token, the exception being a single token of Thai where the boundary falls inside a merged token.

The instruction’s second sentence exists only to make scoring deterministic, so `measure_c.py --no-marker` also reports  $c$  without it, cutting at the last sentence-ending punctuation before the marker: 14/16/19/21 for English, Chinese, Spanish and Vietnamese, 26 for Swahili, 69 for Bengali and 98 for Telugu. Thai joins its two clauses with a conjunction and no punctuation, so it has no separable format sentence and we do not cut it on a guess. The headline cells do not rest on that clause: a minimal one-sentence instruction still leaves Bengali at 69 and Telugu at 98, past the default window, and at  $w=32$  the minimal-instruction blind set is exactly those two.

For a padded or schema-carrying instruction, `measure_c_schema.py` builds the instruction with the same padding routine the item builder uses and reports  $c$  for the padded string, so the measured constant and the generated prompt cannot drift apart.

$Q_{90}$  is produced by `measure_q.py` as the 90th percentile of `len(tokenizer.encode(question))` on a split that excludes the 100 evaluation items: XQuAD validation after index 100, and TyDiQA-GoldP train with every evaluation question id removed. The evaluation items themselves are never used to estimate it.

## C THE SCORING RULE AND ITS ROBUSTNESS

An answer counts as correct if, after NFKC normalization, lowercasing, punctuation stripping and language-aware tokenization (character-level for scripts written without word separators; article stripping only for languages that have articles), some gold span appears inside the marked answer span or inside the first sentence of the prose that precedes the marker. Whole-output containment would over-credit answers that merely quote a passage; marker-only scoring under-credits models that state the answer in prose and then emit a reformatted token after the marker, and that failure rate differs by language, which would manufacture exactly the kind of per-language gap this paper is about. An earlier frozen rule recorded during generation was marker-only; we never quote it and recompute every cell offline.

The choice does not carry the results. Under a stricter marker-only rule the Telugu dose ladder still shows the same shape: the uncompressed baseline scores 50% rather than 56%,  $c+4$  scores 35% and  $c+Q_{90}$  scores 47%, so the hole is  $-15$  pp instead of  $-18$  pp and the  $c+Q_{90}$  cell is  $-3$  pp instead of 0 pp. Applied to the full-pool Telugu re-run of Appendix D, the same rule leaves the certified residual in place:  $-9.4$  pp (25 fixed / 88 broken,  $p=2.1 \times 10^{-9}$ , interval  $[-11.9, -6.5]$ ). Table 3 repeats the 8B Bengali cells under two further scorers.

Table 1 reports 25 paired tests. Under a Holm correction across the grid, five blind cells survive — Bengali at  $w=56$  and 88, Telugu at  $w=32$ , 56 and 88 — and the single-window Thai and Swahili dips, Bengali at  $w=32$  ( $p=.0072$ , Holm threshold .0025), and the Telugu  $w=120/176$  cells do not.

No claim in the paper rests on an unsurviving cell alone: the Telugu  $c+\varepsilon$  point is carried by the dose ladder, whose  $c+4$  cell stands at  $p=.0009$  in a six-test family.

Table 3: Robustness of the 8B Bengali cells to the scoring rule. The spelling fold maps two pairs of Bengali graphemes that the model interchanges; gold-token recall accepts a prediction covering at least 70% of the gold tokens. The ranking and the sign of every comparison are unchanged.

| scorer                            | baseline | $w=64$ | $\hat{w}$ | $\hat{w} - \text{baseline}$ |
|-----------------------------------|----------|--------|-----------|-----------------------------|
| containment (used throughout)     | 81       | 64     | 73        | -8                          |
| containment after a spelling fold | 81       | 65     | 75        | -6                          |
| gold-token recall $\geq 0.7$      | 85       | 70     | 79        | -6                          |

## D CLOSURE, WITH INTERVAL ESTIMATES

The registered closure rule —  $|\Delta| \leq 3$  pp against the uncompressed baseline, as a point estimate — is a decision rule fixed before each run, not an equivalence test, and at  $n=100$  the two are far apart. Table 4 reports an exact 95% interval for every closure cell, obtained by conditioning on the discordant count  $N$ : given  $N$ , the number of items the treatment fixes is binomial, and a Clopper-Pearson interval for its rate maps monotonically onto an interval for  $\Delta$ . Five cells support the stronger reading that the lower bound excludes a loss beyond  $-3$  pp: English at 4B and at Llama, and all three schema tails. The Bengali and Telugu intervals at 4B are consistent with the gate they met but also with losses the gate would have rejected — which is what  $n=100$  can and cannot say — and the same holds for the Gemma, Llama and instruction-first cells. The 8B interval excludes zero, confirming that residual. The recovery claims against the default window ( $+9$  to  $+19$  pp,  $p \leq .022$ ) are unaffected: they are significance claims and carry their own tests.

Two of those cells were then re-run at depth. The stopping rule was fixed before any of it existed: final  $n$  is the entire validation pool, 669 Telugu and 113 Bengali items, a rule with no free parameter and therefore nothing to stop early on. Telugu’s interval now lies entirely below the gate rather than merely below zero — the stronger reading the table calls a certified residual — and the decomposition is exact: the original hundred items reproduce byte-identically and still sit at 0.0, while the 569 new ones sit at  $-6.7$ . Bengali meets the gate on a pool its own size caps at 113, so its interval stays wide, as the preregistration said it would. The 8B Bengali cell is consequently no longer the only residual in this paper that an interval confirms; the primary model now carries one too, measured at  $6.7 \times$  the sample.

## E FULL WINDOW, PAD, AND $c+16$ GRIDS

Table 5 gives the measured constants for all eight languages. Table 6 gives the eight-language run of the  $c+16$  ablation: the five languages whose trailing block already fits inside the default window stay within 3 pp of baseline at both settings, with a marginal Vietnamese exception at the default window; only Bengali and Telugu carry a hole for  $c+16$  to close, and it does not close it. Table 7 gives the English pad sweep plotted in Figure 2b.

## F THE DOSE LADDER IN FULL

The Telugu ladder is the upper block of Table 8. Binning its five treated cells by per-item visibility gives:  $V < 0.25$ ,  $n=118$ ,  $-16.1$  pp (6 fixed / 25 broken,  $p=9 \times 10^{-4}$ );  $0.25 \leq V < 0.5$ ,  $n=104$ ,  $-4.8$  pp;  $0.5 \leq V < 0.75$ ,  $n=84$ ,  $-2.4$  pp;  $0.75 \leq V < 1$ ,  $n=59$ ,  $-11.9$  pp (2/9,  $p=.065$ );  $V=1$ ,  $n=135$ ,  $+0.7$  pp (9/8).

The  $0.75 \leq V < 1$  band is the one that breaks monotonicity, so we describe its contents. It contains nine distinct items with questions of 40 to 81 tokens; six of the nine enter through the  $c+48$  cell alone, and four of the nine are answered correctly once the window reaches  $V=1$ . The band is not significant on its own ( $p=.065$ ) and it represents 59 of the 500 treated observations. Read together with the  $V=1$  bin, it says that visibility close to 1 is not yet equivalent to visibility at 1 for long

Table 4: Exact-conditional 95% intervals for the paired difference at  $\hat{w}$ ; `closure_cis.py`. The instruction-first rows are mechanism-gate cells — the default window with the question last, so  $V=1$  by construction — judged by the same rule.

| cell            | $\Delta pp$ | fixed/broken | 95% CI          | reading                          |
|-----------------|-------------|--------------|-----------------|----------------------------------|
| 4B en           | +2          | 2/0          | $[-1.4, +2.0]$  | non-inferior at $-3$ pp          |
| 4B bn           | -2          | 6/8          | $[-9.1, +5.9]$  | interval too wide for $\pm 3$ pp |
| 4B te           | 0           | 6/6          | $[-6.9, +6.9]$  | interval too wide for $\pm 3$ pp |
| full pool: te   | -5.7        | 19/57        | $[-7.8, -3.1]$  | certified residual               |
| full pool: bn   | -1.8        | 6/8          | $[-8.0, +5.2]$  | gate met; interval wide          |
| JSON-60         | +2          | 3/1          | $[-2.4, +3.9]$  | non-inferior at $-3$ pp          |
| JSON-120        | +1          | 1/0          | $[-0.9, +1.0]$  | non-inferior at $-3$ pp          |
| JSON-200        | +2          | 2/0          | $[-1.4, +2.0]$  | non-inferior at $-3$ pp          |
| Gemma bn        | -6          | 3/9          | $[-10.7, +1.7]$ | interval too wide for $\pm 3$ pp |
| 8B bn           | -8          | 0/8          | $[-8.0, -2.1]$  | confirmed residual               |
| Llama en        | +1          | 1/0          | $[-0.9, +1.0]$  | non-inferior at $-3$ pp          |
| Llama bn        | -8          | 3/11         | $[-12.7, +0.2]$ | interval too wide for $\pm 3$ pp |
| Llama te        | -1          | 3/4          | $[-5.6, +4.4]$  | interval too wide for $\pm 3$ pp |
| IF bn at $w=64$ | -4          | 3/7          | $[-8.7, +3.0]$  | interval too wide for $\pm 3$ pp |
| IF te at $w=64$ | -6          | 1/7          | $[-7.9, +0.4]$  | interval too wide for $\pm 3$ pp |

Table 5: Measured constants, Qwen3-4B.  $I$  is the instruction length,  $c=I+s$  with  $s=5$ , and  $Q_{90}$  is the held-out question-length percentile ( $n=1090/2387/5563$  for en/bn/te;  $Q_{50}$  is 12/46/54 and  $Q_{\max}$  is 36/141/170).  $Q_{90}$  was measured for the languages the AutoWindow arms use and for the two languages the  $w=32$  flip touches in the sweep; Spanish and Vietnamese also have  $c > 32$  but are not swept.

|                      | en | zh | es | vi | th | sw | bn  | te  |
|----------------------|----|----|----|----|----|----|-----|-----|
| $I$                  | 20 | 24 | 30 | 34 | 40 | 42 | 102 | 162 |
| $c$                  | 25 | 29 | 35 | 39 | 45 | 47 | 107 | 167 |
| $c+16$               | 41 | 45 | 51 | 55 | 61 | 63 | 123 | 183 |
| $Q_{90}$             | 18 | —  | —  | —  | 46 | 20 | 76  | 80  |
| $\hat{w} = c+Q_{90}$ | 43 | —  | —  | —  | 91 | 67 | 183 | 247 |

questions — which is the same statement as the registered leftover of a language-level percentile, seen per item.

The Bengali ladder, the lower block, was run on its own grid with the same readout script and repeats the shape. Its bins are  $V < 0.25$ ,  $n=129$ ,  $-14.0$  pp (5 fixed / 23 broken,  $p=9 \times 10^{-4}$ );  $0.25 \leq V < 0.5$ ,  $n=95$ ,  $-7.4$  pp;  $0.5 \leq V < 0.75$ ,  $n=70$ ,  $-5.7$  pp;  $0.75 \leq V < 1$ ,  $n=59$ ,  $-8.5$  pp (2/7, ns);  $V=1$ ,  $n=147$ ,  $0.0$  pp (9/9). Both ends land where the account says they should, and the band that breaks monotonicity is the same one, which makes the caveat above a two-language observation rather than an artefact of one grid. Within items, 18 switchers are correct at high visibility against 1 in the opposite direction ( $p=10^{-4}$ ), the cluster-robust logit on  $V$  gives  $+0.65$  ( $z=2.95$ ,  $p=.003$ ), and AIC prefers  $V$  (637.7) to  $(w-c)$  (639.7) and to  $(w-c) + |Q|$  (641.5) — again a gap of two, which is support rather than a discrimination. An earlier run of this block was cut short at 43 items on its third rung and tied on that last comparison; the tie resolves, narrowly, once the ladder is complete. Its numbers are superseded here, and the 343 generations the two runs share are byte-identical.

## G SCHEMA TAILS: THE MIS-SIZED RUN AND THE RERUN

The first schema run applied the AutoWindow rule but computed  $c$  from the *unpadded* English instruction, giving  $w=41$  — smaller than the default  $w=64$  it was supposed to improve on. It therefore tested a window below the threshold rather than the rule, and the result was correspondingly bad: against uncompressed baselines of 94/95/94 for the 60-, 120- and 200-token schema tails, the default window scored 88/89/89 ( $-6$ ,  $-6^*$ ,  $-5$  pp) and the mis-sized window scored 88/72/75

Table 6: The  $c+16$  ablation across eight languages. Accuracy, with paired  $\Delta pp$  against the uncompressed baseline in parentheses;  $n=100$ ; star: McNemar  $p < .05$ . Shading marks the blind cells, as in Table 1: of the eight languages, the shipped window blinds two.

|    | $c$ | $c+16$ | baseline | $w=64$    | $c+16$  |
|----|-----|--------|----------|-----------|---------|
| en | 25  | 41     | 93       | 93 (+0)   | 94 (+1) |
| zh | 29  | 45     | 83       | 83 (+0)   | 85 (+2) |
| es | 35  | 51     | 88       | 86 (−2)   | 88 (+0) |
| vi | 39  | 55     | 83       | 78 (−5)   | 81 (−2) |
| th | 45  | 61     | 85       | 85 (+0)   | 84 (−1) |
| sw | 47  | 63     | 56       | 53 (−3)   | 56 (+0) |
| bn | 107 | 123    | 73       | 57 (−16*) | 66 (−7) |
| te | 167 | 183    | 56       | 37 (−19*) | 50 (−6) |

Table 7: English pad sweep. Accuracy,  $n=100$  per cell, eviction ratio 0.75. This store carries no uncompressed baseline, so the comparison is across windows; the bold cell in each row is the first window exceeding that pad’s  $c$ , and it is also that row’s maximum.

| pad | $c$ | $w=32$ | 56 | 80        | 104 | 144       |
|-----|-----|--------|----|-----------|-----|-----------|
| 48  | 57  | 81     | 87 | <b>92</b> | 92  | 92        |
| 64  | 73  | 86     | 90 | <b>94</b> | 93  | 93        |
| 96  | 105 | 82     | 90 | 89        | 90  | <b>93</b> |
| 128 | 137 | 78     | 85 | 89        | 89  | <b>91</b> |

Table 8: The dose ladder, on both languages.  $n=100$  per rung, eviction ratio 0.75, paired McNemar against each language’s own uncompressed baseline. Median  $V$  is the median per-item share of the question that falls inside the window. In each block the offsets are the same and the top rung is that language’s own  $c+Q_{90}$ , which is also its shipped  $\hat{w}$ . Two Telugu rungs recur elsewhere:  $w=183$  is its  $c+16$  ablation,  $w=247$  its  $\hat{w}$ . No rung here is blind, so none is shaded.

| $w$                            | slack $w-c$ | median $V$ | accuracy | $\Delta pp$ | fixed/broken | $p$   |
|--------------------------------|-------------|------------|----------|-------------|--------------|-------|
| Telugu, $c=167$ , baseline 56  |             |            |          |             |              |       |
| 171                            | 4           | 0.08       | 38       | −18*        | 5/23         | .0009 |
| 183                            | 16          | 0.31       | 50       | −6          | 5/11         | .21   |
| 199                            | 32          | 0.62       | 52       | −4          | 4/8          | .39   |
| 215                            | 48          | 0.92       | 52       | −4          | 6/10         | .45   |
| 247                            | 80          | 1.00       | 56       | 0           | 6/6          | 1.0   |
| Bengali, $c=107$ , baseline 73 |             |            |          |             |              |       |
| 111                            | 4           | 0.09       | 56       | −17*        | 3/20         | .0005 |
| 123                            | 16          | 0.34       | 66       | −7          | 5/12         | .14   |
| 139                            | 32          | 0.68       | 70       | −3          | 4/7          | .55   |
| 155                            | 48          | 1.00       | 68       | −5          | 3/8          | .23   |
| 183                            | 76          | 1.00       | 71       | −2          | 6/8          | .79   |

(−6, −23\*, −19\* pp). The mis-sizing is itself an instance of the paper’s claim: an English-derived constant applied to a prompt whose trailing block is much longer.

The rerun measures  $c$  with the schema in place (72/166/213) and is reported in panel (B) of Table 2. The two runs share their uncompressed and default-window cells, which the rerun reproduced exactly, generation for generation.

## H A SECOND TOKENIZER: GEMMA-3-4B-IT

Gemma-3-4b-it encodes the same instructions much more compactly than Qwen3-4B:  $c$  is 27/32/44 for English, Bengali and Telugu against Qwen’s 25/107/167, and its questions are shorter in tokens as well ( $Q_{90} = 18/18/20$  against 18/76/80). The threshold should therefore move down with the tokenizer, and the window sweep in Table 9 is consistent with that: Bengali is damaged at the two

windows below its  $c$  and drifts to a non-significant  $-5$  pp at the two above it, while Telugu — whose Qwen counterpart loses 19 points at  $w=64$  — shows no comparable hole anywhere on the grid. English is flat to slightly positive throughout.

Table 9: Gemma-3-4b-it window sweep. Paired  $\Delta$ pp against the Gemma baseline (accuracy 75/62/37 for en/bn/te),  $n=100$ , star: McNemar  $p < .05$ . Shading marks blind cells as in Table 1. Not pooled with any Qwen cell.

|    | $c$ | $w=16$ | 24   | 32  | 48 | 64 |
|----|-----|--------|------|-----|----|----|
| en | 27  | +2     | +1   | +3  | +4 | +3 |
| bn | 32  | -13*   | -10* | -9* | -5 | -5 |
| te | 44  | -3     | -4   | -3  | -1 | -3 |

Three qualifications belong with this table. Bengali at  $w=32$ , which is  $c$  exactly, is still  $-9$  pp, so the step is not a clean discontinuity at  $c$  here either; Telugu’s baseline of 37% is low enough that the sweep would not resolve a small hole; and English is itself blind at  $w=16$  and 24 on this tokenizer without being damaged there — indeed no window we tested damages Gemma’s English at this ratio. Blindness is necessary for the failure but not sufficient. How much any ranking is worth is measured in Appendix K: evicting at the same ratio with no ranking at all collapses generation itself, so even a blinded ranking does most of the work. What distinguishes Gemma’s undamaged blind English cells from Qwen’s damaged padded ones, that measurement cannot say, and we report the asymmetry as an observation rather than an account. The within-model contrast is untouched — the same two blind windows cost Bengali 13 and 10 points — but the shading makes the asymmetry visible, and we would rather state it than let a reader find it. The AutoWindow arm on this tokenizer gives  $\hat{w} = 45/50/64$  and returns  $+3/-6/-3$  pp against baseline, against  $+3/-5/-3$  at the default window. For Telugu the rule selects the default constant itself. For Bengali it does not close, and the main text explains why we read that residual as a plateau rather than a threshold effect: it is present at every window from 48 upward, and at  $\hat{w}=50$  the median Bengali question of 11 tokens is entirely inside the window.

## I SCALE: QWEN3-8B

The scale slice covers English and Bengali with the same formula and the same 100 items per cell. English is at ceiling and carries no information: 96 baseline, 97 at the default window, 96 at  $\hat{w}=43$ . Bengali reproduces the phenomenon: 81 baseline, 64 at the default window ( $-17$  pp, 1 fixed / 18 broken,  $p<.001$ ), and 73 at  $\hat{w}=183$ , which recovers  $+9$  pp against the default window (11/2,  $p=.022$ ) but stays 8 pp below baseline (0/8,  $p=.008$ ). The registered criterion was  $\pm 3$  pp, so this is a miss, and we report it as one.

What the residual is not is worth stating precisely, because it is the same diagnosis that applies to Gemma. Of the eight items that the baseline answers and  $\hat{w}$  does not, all eight have questions shorter than  $Q_{90}$  (38 to 67 tokens), so the window covers the question entirely for every one of them; seven of the eight reproduce the question verbatim in the generated text; and all twelve items whose question exceeds  $Q_{90}$  — the only items where the language-level percentile could still leave the scorer blind — are answered correctly under  $\hat{w}$ . Among the items that the default window breaks and  $\hat{w}$  does not recover, three of seven run to the decode cap without ever emitting an answer marker, against none of the eleven items it does recover, and six of the seven sit at the middle gold position. The errors that remain are near-copies of the gold span (a dropped morpheme, an interchanged grapheme) or answers whose supporting sentence did not survive eviction — ordinary compression damage, not blindness.

## J A THIRD FAMILY: LLAMA-3.1-8B-INSTRUCT

Both inputs to the rule are re-measured on the Llama tokenizer, which is costlier than Qwen’s on both: the trailing blocks are 25/125/190 tokens for English, Bengali and Telugu (Qwen: 25/107/167) and the held-out question percentiles are  $Q_{90}=18/87/94$  (Qwen: 18/76/80), giving  $\hat{w} = 43/212/284$ . Baselines are 97/76/51 (Table 10).



Bengali reproduces the hole at the default window —  $76 \rightarrow 58$ ,  $-18$  pp (2 fixed / 20 broken,  $p=.0001$ ), with a 95% interval of  $[-21.5, -9.2]$ , the only hole in the paper whose interval sits entirely below the  $-8$  damage bound — and  $\hat{w}=212$  recovers  $+10$  pp (13/3,  $p=.021$ ) while remaining  $-8$  pp below baseline ( $p=.057$ ). English is at ceiling (97/98/98). Telugu shows no clear hole at the default ( $-5$ , ns) and  $\hat{w}$  meets the gate ( $-1$ ), but the cell is low-confidence: 88 of 100 Telugu outputs never emit the answer marker, so scoring rests entirely on the first-sentence rule and the marker-only robustness check that backs the Qwen tables does not exist here. We quote Telugu-Llama only as “no clear hole; the formula does not hurt”.

The residual audit is more mixed than at 8B. Of the eleven items that the baseline answers and  $\hat{w}$  does not, nine have questions within  $Q_{90}$  (the window covers them fully; seven of the eleven sit at the middle gold position, and 4 of 10 unrecovered default-window breaks ramble to the decode cap), while the other two belong to the long decile — the first arm in which the registered percentile leftover is actually visible: 10 of 12 long-question items are correct under  $\hat{w}$ , against 12 of 12 at 8B.

Table 10: Llama-3.1-8B-Instruct,  $n=100$  per cell, eviction ratio 0.75, own stack. Accuracy, with the paired difference against this model’s uncompressed baseline and its exact 95% interval.

|    | $c$ | $\hat{w}$ | baseline | $w=64$         | at $\hat{w}$                         |
|----|-----|-----------|----------|----------------|--------------------------------------|
| en | 25  | 43        | 97       | 98 (+1)        | <b>98</b> (+1; $[-0.9, +1.0]$ )      |
| bn | 125 | 212       | 76       | 58 ( $-18^*$ ) | <b>68</b> ( $-8$ ; $[-12.7, +0.2]$ ) |
| te | 190 | 284       | 51       | 46 ( $-5$ )    | <b>50</b> ( $-1$ ; $[-5.6, +4.4]$ )  |

## K RIVAL REMEDIES, AND A SECOND MEMBER OF THE FAMILY

Two remedies compete with sizing the window: switch to a scorer that never reads the question, so that nothing can be hidden from it, or ship one larger constant and stop measuring. Both were run on the same items, at matched ratios, against each store’s own uncompressed baseline (Table 11). A third block asks the complementary question: whether the integer we compute is a property of the prompt or of the method it was first tested on.

**Scorers that do not read the question.** Expected Attention ranks by an estimate of future attention (Devoto et al., 2025) and TOVA by attention already paid (Oren et al., 2024); neither can be blinded by a trailing block, because neither was looking. At  $r=0.75$  both lose significantly in every language — including English, where the trailing block is 25 tokens and no window we tested costs anything. Their damage does not follow  $c$ , which is what distinguishes it from the failure this paper describes; it is simply the price of ranking a cache without knowing what is being asked. Blindness is a property of a window that can be sized; query-agnosticism is a property of the method.

One prediction registered for this arm cannot be scored as written: it compares Thai against Bengali, and the arm as locked runs English, Bengali and Telugu. We flag the drafting error rather than substitute a different comparison after the fact. The claim it was meant to test — that this damage does not track  $c$  — is carried by the English column, which was written into the same preregistration.

**One larger constant.** At  $r=0.9375$  the retained cache is about 512 tokens on these 8k prompts, so a  $w=256$  window protects half of what survives. It costs English 9 pp against baseline ( $p=.049$ ; against the default window the same gap is  $-7$  and does not reach significance at  $n=100$ ), while  $\hat{w}_{\text{en}}=43$  costs one. It does not beat  $\hat{w}$  anywhere:  $-8$  pp on English ( $p=.077$ ) and  $-4$  pp on Bengali (ns), so the dominance is carried by English and we do not claim it per language. Bengali collapses at the default window here ( $-44$  pp) and  $\hat{w}$  recovers  $+20$  pp of that (22 fixed / 2 broken,  $p<10^{-4}$ ), its largest recovery at this ratio — and still finishes 24 pp below baseline. Sizing the window is what removes the blind mode; it is not what makes aggressive eviction free.

The reviewer-facing question is the other ratio: does the constant suffice where the headline numbers live? Measured on its own pod and baselines, judged by the same preregistered  $\pm 3$  gate the rule was judged by,  $w=256$  at  $r=0.75$  is  $+1$  on English,  $-3$  on Telugu — at the boundary — and  $-5$  on Bengali (5 fixed / 10 broken, ns, CI  $[-11.5, +3.5]$ ), missing the gate that  $\hat{w}$  met at  $-2$ . The preregistration expected the constant to pass and committed to rewriting our claim if it did; it did

not, and we state the symmetric caveat instead: at  $n=100$  the Bengali interval certifies neither the miss nor a close, so what the grid supports is that the rule met all six of its gates and the constant met four of six across the two ratios, never beating the rule’s cell anywhere.

**No ranking at all.** The same store carries the control the registry always reserved: `random@r0.75`, eviction with no scorer. It does not lose answers; it collapses generation — 6/1/0% accuracy on English, Bengali and Telugu, with 96–100% of outputs running to the decode cap, mostly as repetition loops. Even a blinded ranking is carrying enormous weight. Because random sits at the floor for every language, it bounds the value of ranking from above without discriminating between them, and we draw no per-language conclusion from it.

**A second member of the family.** PyramidKV (Cai et al., 2024) keeps the observation window and changes what is done with the retained budget, allocating more of it to lower layers; its `kvpress` implementation subclasses `SnapKV` and inherits  $w=64$  unchanged. On its own pod and its own baselines, that inherited default digs deeper holes than `SnapKV`’s —  $-28$  pp on Bengali and  $-32$  on Telugu, against  $-16$  and  $-19$  — and the `AutoWindow` integers, used exactly as shipped, remove them:  $-2$ ,  $-1$  and  $+0$  on Bengali, Telugu and English, with recoveries of  $+26$  (29 fixed / 3 broken) and  $+31$  (32/1), the two largest in this paper. We registered the opposite guess about the depth, expecting the pyramid allocation to shallow the blind hole rather than deepen it, and report the outcome as an observation we did not predict. The transfer itself is the point: an integer computed from a tokenizer and a template, never tuned against accuracy on either method, repairs both.

Table 11: Rival remedies, paired  $\Delta$ pp against each store’s own uncompressed baseline,  $n=100$  per cell; star: McNemar  $p < .05$ . Top: two query-agnostic scorers at the ratio used throughout the paper, beside the windowed cells they would replace. Bottom: the heavy-ratio sweep, where  $\hat{w}$  is 43 for English and 183 for Bengali. Third block: the constant and the no-ranking control at the headline ratio; fourth: PyramidKV. Both ran on pods that did not reproduce the campaign stack, so their bracketed `SnapKV` values are context rather than comparators and are never pooled with them.

| $r=0.75$           | Exp. Attention | TOVA                   | <code>SnapKV</code> $w=64$                 | <code>SnapKV</code> $\hat{w}$ |
|--------------------|----------------|------------------------|--------------------------------------------|-------------------------------|
| en                 | $-14^*$        | $-8^*$                 | $+0$                                       | <b><math>+2</math></b>        |
| bn                 | $-18^*$        | $-16^*$                | $-16^*$                                    | <b><math>-2</math></b>        |
| te                 | $-20^*$        | $-18^*$                | $-19^*$                                    | <b><math>+0</math></b>        |
| $r=0.9375$         | $w=64$         | $w=256$                | $\hat{w}$                                  |                               |
| en                 | $-2$           | $-9^*$                 | $-1$                                       |                               |
| bn                 | $-44^*$        | $-28^*$                | <b><math>-24^*</math></b>                  |                               |
| $r=0.75$ , own pod | random         | $w=256$                | [ <code>SnapKV</code> $\hat{w}$ ]          |                               |
| en                 | $-87^*$        | <b><math>+1</math></b> | [ $+2$ ]                                   |                               |
| bn                 | $-72^*$        | $-5$                   | [ $-2$ ]                                   |                               |
| te                 | $-56^*$        | <b><math>-3</math></b> | [ $+0$ ]                                   |                               |
| PyramidKV          | $w=64$         | $\hat{w}$              | [ <code>SnapKV</code> : $w=64 / \hat{w}$ ] |                               |
| en                 | $+2$           | <b><math>+0</math></b> | [ $+0 / +2$ ]                              |                               |
| bn                 | $-28^*$        | <b><math>-2</math></b> | [ $-16^* / -2$ ]                           |                               |
| te                 | $-32^*$        | $-1$                   | [ $-19^* / +0$ ]                           |                               |

## L A SECOND CONSTANT: THE DECODE CAP

The main arms decode under a 384-token cap, and Appendix N says why in a sentence; this appendix is the measurement behind it. The campaign’s first arms used 128, and a decode cap is itself a token-denominated constant — so the paper’s thesis applies to it, one layer further out, in the evaluation harness rather than in the evictor. The evidence comes from the cap-128-era stores: a different stack and the byte-parallel item variant, whose inputs are byte-equal across languages. Nothing here is pooled with the main tables, and every number is recomputed by `decode_cap_ledger.py`.

At a 128-token cap, the share of *uncompressed* baseline generations that hit the cap runs from 1% in Chinese to 65% in Greek (Table 12) — a spread produced by one integer, on byte-equal inputs,

with no compression anywhere in the picture. Because the instruction asks for the answer marker at the end of the reply, truncation does not merely shorten an answer: it decides which branch of the scoring rule fires. At 128, 95% of English baselines carry the marker against 45% of Greek; at 384 the marker comes back (90/97/92% for Greek, Hindi and Russian). An English-validated harness would see none of this.

Greedy decoding makes the 128-token output a byte prefix of the 384-token output on 1,124 of the 1,200 generations the two runs share, so the cap’s effect is isolable within an item rather than inferred across conditions. On the 327 generations where truncation destroys the marker, our scoring rule loses 4 pp (19 fixed / 7 broken,  $p=.029$ ; Greek +2 and Hindi +1, both ns, Russian +11,  $p=.039$ ). That is real and it is bounded, because the rule’s prose-fallback branch was built to absorb exactly this failure. We cannot say from these stores what a format-strict harness would lose — our own strict marker path also fails for reformatting reasons that have nothing to do with the cap — so we claim only the part we can measure: the cap selects the scoring branch per language, and under our rule the residual cost is four points on the cells it touches. This is why the main arms decode at 384, and why we treat every token-denominated constant in the pipeline — window, cap, budget — as a quantity to measure rather than a default to inherit.

Table 12: The decode cap as a second token-denominated constant. Uncompressed baselines on byte-parallel items,  $n=100$  per language, from the cap-128-era stores; Greek, Hindi and Russian were re-run at 384. Truncated: share of generations that reach the cap. Marker: share carrying the end-of-reply answer marker that the scoring rule keys on.

|                  | el  | hi  | ru  | th  | de  | en  | vi  | es  | zh  |
|------------------|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| truncated at 128 | 65% | 50% | 30% | 14% | 12% | 6%  | 4%  | 4%  | 1%  |
| marker at 128    | 45% | 59% | 71% | 58% | 74% | 95% | 94% | 96% | 98% |
| truncated at 384 | 17% | 4%  | 2%  | —   | —   | —   | —   | —   | —   |
| marker at 384    | 90% | 97% | 92% | —   | —   | —   | —   | —   | —   |

## M FROZEN INSTRUCTIONS AND PROMPT LAYOUT

The QA instructions are one line per language, frozen under a prompt version key so that a change to them would invalidate the stored runs rather than silently alter them. They carry the same communicative content — answer from the passages above, end with a marked span — and differ only in how many tokens the model’s tokenizer needs for it: 20 for English, 102 for Bengali, 162 for Telugu on Qwen3-4B. Every instruction-last item has the layout

```
[gold + distractor passages, packed to 8k tokens]
<question>
<instruction>
<chat suffix>          # Qwen: <|im_end|>\n<|im_start|>assistant\n
```

and the observation window is the last  $w$  tokens of that concatenation. In English, for instance:

```
... Paris is the capital of France. ...
What is the capital of France?
Answer the question using only the passages above.
End your reply with '#### <exact answer span>'.
<|im_end|>
<|im_start|>assistant
```

The padded conditions prepend content-free filler to the instruction so that the marked-span sentence stays adjacent to generation and only the token count changes; the schema conditions prepend a JSON response specification in the same position.

## N MEASUREMENT TRANSCRIPTS AND RUN PROVENANCE

**Stacks.** Each generation record stores a stack descriptor (GPU, driver, CUDA, torch, transformers, kvpress) and every comparison in this paper is within one descriptor. The Gemma cells share one stack and the Qwen3-8B slice another; the Qwen3-4B cells share a third, with one exception — the PyramidKV and constant-control pods each landed on descriptors of their own (a newer driver patch or a different host kernel; same GPU model and library versions), so those arms pair only within themselves and are quoted beside the others rather than pooled with them. We never pool across descriptors.

**Determinism, measured.** The follow-up runs were executed weeks after the original ones, on new machines, and they re-ran cells that the earlier stores already contained — uncompressed baselines, default-window cells, the rungs a later ladder shares with an earlier one. We do not count these by hand: `determinism_ledger.py` keys every generation by model, task, language, treatment, item and prompt version, takes the earliest run of each key as the original, and compares every later run against it. It finds 3,843 generations produced a second time in a later store, and all 3,843 reproduce the original text byte for byte. This is what allows us to treat the re-measured constants and the earlier tables as describing the same system. Three hundred of those pairs also cross a naming difference — one store requests the default window implicitly and the other names  $w=64$  explicitly — and agree token for token, as they should. The two drifted pods bought a measurement we had not planned: the 600 uncompressed baselines they share with the AutoWindow store are byte-identical *across* those boundaries. The ledger counts same-descriptor repeats (3,843, the figure above) and cross-descriptor ones (600) separately, and the argument that the earlier and later tables describe one system rests only on the former.

**Decode cap.** Every arm reported in the main text uses a 384-token decode cap. The supporting window-widening comparison in Appendix O was measured under a 128-token cap, as were the historical instruction-first numbers kept beside it; the instruction-first layout swap itself has been re-measured at the current cap and is reported in the main text (Table 2). We flag the difference here rather than leave it to be discovered in a log. The cap matters, and it is an instance of this paper’s subject rather than a detail of it: at 128 tokens the share of answers truncated varies from 1 to 65% across languages, which is why the main arms were re-run at 384 (Appendix L).

## O TWO FURTHER LEVERS ON VISIBILITY

Two earlier interventions move the same quantity and are consistent with the account in the main text. Both were measured under the 128-token decode cap noted above, so we report them as supporting rather than as headline evidence; the layout swap has since been re-measured at the current cap and moved to the main text. The earlier cap-128 measurement is retained here only as provenance: Bengali +14 pp at a fixed ratio ( $34/6$ ,  $p < 10^{-5}$ ,  $n=200$ ) and +22 pp under an absolute cache budget ( $51/7$ ,  $p < 10^{-8}$ ,  $n=200$ ), with a Thai cost of 6.5 pp. The current-cap rerun locates that Thai cost in the uncompressed baseline (a layout effect, −5 pp, ns) rather than in compression, and notes that Thai drops the answer marker in 43% of question-last outputs, so part of the shift is format behaviour. Widening the window from 64 to 128 tokens at a fixed ratio recovers Bengali by 10.5 pp (22 fixed / 1 broken,  $p < 10^{-5}$ ,  $n=200$ ) and saturates there: going on to 256 tokens changes nothing (−0.5 pp, 8/9). Saturation at a window well below  $c + Q_{90}$  for that setting is consistent with the dose curve in Section 5.

## P THE LAYOUT IN THE WILD

Our items are constructed, which invites a fair question: does any widely used prompt actually put a long block after the question? We surveyed 22 templates selected by popularity and officialness *before* their bodies were read — the default QA prompt of a framework above 5k stars, an official vendor documentation example, or the prompt of a published benchmark or harness — so that safe layouts could not be filtered out. Each was recorded verbatim from a commit-pinned source, the list was frozen, and only then was anything tokenized. We measure the *static* trailing block: the

text after the question placeholder with every unfilled placeholder deleted, on the Qwen3-4B tokenizer. Deleting placeholders makes every number a lower bound, because runtime content only adds (Table 13).

This is a measurement of geometry, not of damage. Whether windowed eviction harms an agent loop the way it harms instruction-last QA is not tested here, and the question an agent is answering at step  $k$  is not only the task it was given at step 0. What the survey supports is narrower: the layout in which a token-denominated window can blind the scorer is not an artefact of our construction.

Most of what ships is safe, and that is the first half of the result. Single-shot retrieval-QA defaults leave 4–5 tokens after the question, and their chat variants leave none. Modern vendor structured-output and tool-use APIs carry the schema in an API field rather than in prompt text, so they leave nothing after the task at all — which falsifies the guess we began with, that in-prompt response schemas would be where long trailing blocks live. Harness QA templates leave 3–5 tokens.

The blocks appear where a prompt is *iterated*. LlamaIndex’s shipped refine template puts the query first and, at its empty-placeholder floor, exactly 64 tokens after it — the kvpress default, to the token, before the existing answer or the new context contributes anything. SWE-agent’s default instance template puts 221 static tokens after the problem statement, longer than the longest trailing block among our eight languages, before its first observation arrives. ReAct and the OpenHands loop render the user’s question once and append every later tool result after it by construction. The full record, with permalinks and verbatim template text, ships with the code.

Table 13: Static trailing block of 22 popular templates, in tokens on the Qwen3-4B tokenizer, with unfilled placeholders deleted so each number is a lower bound. Zero means the question is last. Selection was frozen before any template body was read or tokenized.

|                                       | template                           | tokens     |
|---------------------------------------|------------------------------------|------------|
| <i>Single-shot retrieval QA</i>       |                                    |            |
| T01                                   | LangChain RetrievalQA, completion  | 5          |
| T02                                   | LangChain RetrievalQA, chat        | 0          |
| T03                                   | LangChain hub retrieval-qa-chat    | 0          |
| T04                                   | LlamaIndex text-QA default         | 4          |
| T05                                   | LlamaIndex text-QA, chat           | 4          |
| T08                                   | Haystack PromptBuilder example     | 5          |
| T09                                   | RAGFlow chat assembly              | 0          |
| <i>Structured output and tool use</i> |                                    |            |
| T10                                   | OpenAI structured outputs          | 0          |
| T11                                   | OpenAI function calling            | 0          |
| T12                                   | OpenAI JSON mode example           | 6          |
| T13                                   | Anthropic tool use                 | 0          |
| T14–T16                               | Instructor, Outlines, Guidance     | 0          |
| <i>Evaluation harnesses</i>           |                                    |            |
| T20                                   | lm-eval XQuAD, English             | 3          |
| T21                                   | lm-eval XQuAD, Arabic              | 5          |
| T22                                   | Belebele, official / harness       | 19 / 15    |
| <i>Iterated prompts</i>               |                                    |            |
| T07                                   | LlamaIndex refine, chat            | 9          |
| T17                                   | LangChain ReAct, static suffix     | 3          |
| T06                                   | <b>LlamaIndex refine default</b>   | <b>64</b>  |
| T18                                   | <b>SWE-agent instance template</b> | <b>221</b> |
| T19                                   | OpenHands mid-loop                 | grows      |

## Q PRODUCTION SCOPE

Re-checked on 2026-08-18, with permalinks in the supplementary record. vLLM exposes KV-cache dtype and scheduler token budgets, not a SnapKV-style evictor; its default attention backends are dense. TensorRT-LLM implements last- $w$  scoring as an opt-in sparse-attention backend whose configuration default is a 32-token window, and skips it entirely below a prompt-length threshold.

SGLang has an open proposal for a sparsity framework which states that it must be default-off and names SnapKV and PyramidKV among its explicit non-goals — so the family we study is, for now, deliberately outside the default path of the two most used open serving stacks. Mao et al. (2026) argue that research evictors rarely ship as production defaults, and we cite that against ourselves: the constant we study is a research default and one opt-in production constant, not something most deployments are running today. The reason to fix it now is that the fix is a measurement rather than a redesign.